

## **Analysis of 10 AI Model Providers**

### **1. OpenAI**

- **Model Name:** GPT-5.5
- **Release Date:** April 2026
- **Knowledge Cutoff Date:** **December 01, 2025**
- **Cost of Token:** 0.50 / \$30.00 per 1M output tokens. (*Prompts exceeding 272K input tokens are priced at 2x input and 1.5x output*).
- **Context Window:** 1,050,000 tokens (with 128,000 max output tokens)
- **Public or Private:** Public (accessible via API Chat Completions, Responses API, etc.)
- **When to Use:** Billed as "a new class of intelligence for coding and professional work," it is ideal for the most complex professional tasks. It features highly advanced reasoning (with configurable reasoning effort levels) and supports comprehensive tool usage including Web Search, File Search, Code Interpreter, Hosted Shell, and Computer Use.
- **Key Information:** GPT-5.5 officially supports Text and Image input, outputting Text. It heavily emphasizes agentic features and tool integration rather than native audio/video modalities.

### **2. Anthropic**

- **Model Name:** Claude 4.7 Opus
- **Release Date:** January 2026
- **Knowledge Cutoff Date:** January 2026
- **Cost of Token:** ~75.00 per 1M output tokens
- **Context Window:** 200,000 tokens
- **Public or Private:** Public (accessible via API and Claude Pro)
- **When to Use:** Highly complex reasoning, coding architecture design, deep research, and tasks requiring meticulous safety protocols.
- **Key Information:** Claude 4.7 Opus is Anthropic's most capable frontier model, offering state-of-the-art benchmark performance. It is particularly favored in academic and scientific research for its nuanced cognitive processing.

### **3. Google DeepMind**

- **Model Name:** Gemini 2.5 Pro
- **Release Date:** Late 2025
- **Knowledge Cutoff Date:** January 2025 , Real-time (Bypasses fixed cutoff via native Google Search integration)
- **Cost of Token:** ~10.00 per 1M output tokens
- **Context Window:** 1,000,000 to 2,000,000 tokens
- **Public or Private:** Public (accessible via Google Cloud Vertex AI and Gemini Advanced)

- **When to Use:** Multimodal processing (native video, audio, and image analysis), massive document processing, and tasks requiring up-to-the-minute web data.
- **Key Information:** Gemini 2.5 Pro boasts one of the largest reliable context windows in the industry, making it the premier choice for parsing entire codebases or hours of raw video and audio feeds in a single prompt.

#### 4. Meta

- **Model Name:** Llama 4 (Maverick)
- **Release Date:** Late 2025
- **Knowledge Cutoff Date:** August 2025
- **Cost of Token:** Varies by cloud provider (Open-weights; hosting costs apply, often <\$0.50 per 1M tokens)
- **Context Window:** 128,000 tokens
- **Public or Private:** Public (Open-weights for commercial and research use)
- **When to Use:** Cost-effective deployment, on-premise hosting, fine-tuning for proprietary enterprise applications, and local software development.
- **Key Information:** Continuing Meta's commitment to the open-source community, Llama 4 matches proprietary models in performance while allowing developers full control over their infrastructure and data privacy.

#### 5. xAI

- **Model Name:** Grok 4
- **Release Date:** Early 2026
- **Knowledge Cutoff Date:** Early 2026 (Supplemented by real-time X/Twitter data)
- **Cost of Token:** Est. 15.00 per 1M output
- **Context Window:** 256,000 tokens
- **Public or Private:** Public (accessible via X Premium+ and xAI API)
- **When to Use:** Real-time social media sentiment analysis, unfiltered brainstorming, and coding (via Grok Code Fast integration).
- **Key Information:** Grok 4 distinguishes itself with unfiltered, real-time access to the X platform's data stream, providing it with a unique advantage in understanding current events, internet culture, and immediate news.

#### 6. Mistral AI

- **Model Name:** Mistral Medium 3
- **Release Date:** May 2025
- **Knowledge Cutoff Date:** Early 2025
- **Cost of Token:** ~6.00 per 1M output
- **Context Window:** 128,000 tokens
- **Public or Private:** Public (accessible via API and cloud platforms)
- **When to Use:** European language translation, cost-effective enterprise scaling, and structured data generation (JSON/SQL).

- **Key Information:** Mistral Medium 3 offers one of the best price-to-performance ratios in the market. It is highly optimized for multi-lingual tasks, particularly European languages, and is designed for highly efficient inference.

## 7. DeepSeek

- **Model Name:** DeepSeek-V3
- **Release Date:** Late 2025
- **Knowledge Cutoff Date:** Late 2024
- **Cost of Token:** Highly subsidized / Open-weights (hosting costs apply)
- **Context Window:** 128,000 tokens
- **Public or Private:** Public (Open-weights)
- **When to Use:** Mathematics, heavy algorithmic programming, and budget-conscious academic research.
- **Key Information:** DeepSeek-V3 shocked the industry by achieving frontier-level math and coding capabilities with a fraction of the training compute used by western competitors, making it a favorite in the developer community.

## 8. Microsoft

- **Model Name:** Phi-4
- **Release Date:** December 2025
- **Knowledge Cutoff Date:** Late 2025
- **Cost of Token:** Varies by Azure deployment (Extremely low latency/cost)
- **Context Window:** 128,000 tokens
- **Public or Private:** Public (Open-weights/Azure)
- **When to Use:** On-device AI (edge computing), mobile applications, and highly specific fine-tuned tasks where massive parameter sizes are unnecessary.
- **Key Information:** Phi-4 is a Small Language Model (SLM). Despite its small parameter count, it utilizes high-quality "textbook" training data to punch well above its weight, matching larger models in reasoning tasks while running efficiently on local hardware.

## 9. Alibaba Cloud (Qwen)

- **Model Name:** Qwen3 Max
- **Release Date:** Early 2026
- **Knowledge Cutoff Date:** Late 2025
- **Cost of Token:** ~6.00 per 1M output
- **Context Window:** 256,000 tokens
- **Public or Private:** Public (accessible via API and HuggingFace for smaller variants)
- **When to Use:** Bilingual (English/Chinese) applications, Asian market localization, and multimodal vision-language tasks.
- **Key Information:** Qwen3 Max is Alibaba's flagship proprietary model. It dominates benchmarks in Chinese language processing while remaining highly competitive with western models in English logic and programming.

## 10. Magic (Magic AI)

- **Model Name:** Magic LTM-2-Mini
- **Release Date:** January 2026
- **Knowledge Cutoff Date:** Late 2025
- **Cost of Token:** Pricing restricted to enterprise contracts
- **Context Window:** 100,000,000 tokens
- **Public or Private:** Private / Limited Enterprise Access
- **When to Use:** Entire software repository ingestion, autonomous software engineering, and continuous long-term memory tasks.
- **Key Information:** Magic's LTM (Long-Term Memory) models feature a revolutionary architecture designed specifically for software engineers. With a groundbreaking 100-million token context window, it can hold entire enterprise codebases in memory simultaneously, allowing for unprecedented code generation and debugging context.