

Analysis of 10 AI Model Providers

1. OpenAI

- **Model Name:** GPT-5.5 Pro
- **Release Date:** 23 Apr 2026
- **Knowledge Cutoff Date:** 01 Dec 2025
- **Cost of Token:** 180.00 per 1M output tokens. (Prompts exceeding 272K input tokens are priced at 2x input and 1.5x output).
- **Context Window:** 1,050,000 tokens (with 128,000 max output tokens)
- **Public or Private:** Public (accessible via API Chat Completions, Responses API, etc.)
- **When to Use:** Problem solving, coding. Best overall reasoning and long-context agents.

2. Anthropic

- **Model Name:** Claude Opus 4.7
- **Release Date:** 16 Apr 2026
- **Knowledge Cutoff Date:** 01 Jan 2026
- **Cost of Token:** 25.00 per 1M output tokens
- **Context Window:** 1,000,000 tokens
- **Public or Private:** Public (accessible via API and Claude Pro)
- **When to Use:** Strongest coding and complex multi-step reasoning.

3. Google DeepMind

- **Model Name:** Gemini 3.1 Pro
- **Release Date:** 19 Feb 2026
- **Knowledge Cutoff Date:** 01 Jan 2025 (can bypass fixed cutoff via native Google Search integration for real-time data)
- **Cost of Token:** 12.00 per 1M output tokens
- **Context Window:** 1,048,576 tokens
- **Public or Private:** Public (accessible via Google Cloud Vertex AI and Gemini Advanced)
- **When to Use:** Reasoning, massive document processing, and native video/audio/image analysis.

4. Meta

- **Model Name:** Llama 4 (Maverick)
- **Release Date:** Late 2025
- **Knowledge Cutoff Date:** August 2025
- **Cost of Token:** Varies by cloud provider (Open-weights; hosting costs apply, often <\$0.50 per 1M tokens)

- **Context Window:** 128,000 tokens
- **Public or Private:** Public (Open-weights for commercial and research use)
- **When to Use:** Cost-effective deployment, on-premise hosting, fine-tuning for proprietary enterprise applications, and local software development.

5. xAI

- **Model Name:** Grok 4.3
- **Release Date:** 30 Apr 2026
- **Knowledge Cutoff Date:** 01 Dec 2025 (Supplemented by real-time X/Twitter data)
- **Cost of Token:** 2.50 per 1M output tokens
- **Context Window:** 1,000,000 tokens
- **Public or Private:** Public (accessible via X Premium+ and xAI API)
- **When to Use:** social media sentiment analysis and unfiltered brainstorming.

6. Mistral AI

- **Model Name:** Mistral Medium 3
- **Release Date:** May 2025
- **Knowledge Cutoff Date:** Early 2025
- **Cost of Token:** ~6.00 per 1M output tokens
- **Context Window:** 128,000 tokens
- **Public or Private:** Public (accessible via API and cloud platforms)
- **When to Use:** European language translation, cost-effective enterprise scaling, and structured data generation (JSON/SQL).

7. DeepSeek

- **Model Name:** DeepSeek V4 Pro
- **Release Date:** 24 Apr 2026
- **Knowledge Cutoff Date:** 01 May 2025
- **Cost of Token:** 0.87 per 1M output tokens
- **Context Window:** 1,048,576 tokens
- **Public or Private:** Public (Open-weights / API)
- **When to Use:** Frontier-level reasoning, complex coding, and long-context agent workflows.

8. Microsoft

- **Model Name:** Phi-4
- **Release Date:** December 2025
- **Knowledge Cutoff Date:** Late 2025
- **Cost of Token:** Varies by Azure deployment (Extremely low latency/cost)
- **Context Window:** 128,000 tokens

- **Public or Private:** Public (Open-weights/Azure)
- **When to Use:** On-device AI (edge computing), mobile applications, and highly specific fine-tuned tasks where massive parameter sizes are unnecessary.

9. Alibaba Cloud (Qwen)

- **Model Name:** Qwen 3.6 Max-Preview
- **Release Date:** 20 Apr 2026
- **Knowledge Cutoff Date:** 01 Mar 2025
- **Cost of Token:** ~0.28 per 1M output tokens
- **Context Window:** 1,000,000 tokens
- **Public or Private:** Public (accessible via API)
- **When to Use:** Frontier agentic coding and benchmark-heavy tasks. Excellent for bilingual (English/Chinese) applications.

10. Magic (Magic AI)

- **Model Name:** Magic LTM-2-Mini
- **Release Date:** January 2026
- **Knowledge Cutoff Date:** Late 2025
- **Cost of Token:** Pricing restricted to enterprise contracts
- **Context Window:** 100,000,000 tokens
- **Public or Private:** Private / Limited Enterprise Access
- **When to Use:** Entire software repository ingestion, autonomous software engineering, and continuous long-term memory tasks.